

Naive Bayes predictor

```
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import StratifiedKFold
import numpy as np
import pickle
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

Load the Data Set

will be split to test and train

```
X = np.load("../X.npy")
y = np.load("../y.npy")

with open("../label_map.pkl", "rb") as f:
    label_map = pickle.load(f)
```

Split the data to n-fold

```
nfold = StratifiedKFold(n_splits=5, shuffle=True, random_state=777)
fold = 1
```

Separate them into 5 different modules

to see how the module will act on the data

```
from sklearn.metrics import precision_recall_fscore_support

accuracies = []
precisions, recalls, f1s = [], [], []
conf_matrices = []
scaler = StandardScaler()

for training_index, testing_index in nfold.split(X, y):# the loop calls next()
    automatically to give the proper splitted sets for the fold
    # train_index: A NumPy array holds all indices of the training items in
    this fold, e.x. [0 4 7 ... 2887 2996 2999]

    # Get the training and testing data for this fold
    X_train, X_test = X[training_index], X[testing_index]
    y_train, y_test = y[training_index], y[testing_index]

    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)
```

```

# Normalize feature values
"""
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
"""

# Train Naive Bayes model
model = GaussianNB()
model.fit(X_train, y_train)

# Predict and evaluate
y_pred = model.predict(X_test)

# Save performance measurements
# Accuracy
acc = accuracy_score(y_test, y_pred)
accuracies.append(acc)
# Precision, Recall, F1 (macro average)
p, r, f, _ = precision_recall_fscore_support(y_test, y_pred,
average='macro')
precisions.append(p)
recalls.append(r)
f1s.append(f)
# Confusion matrix (optional)
cm = confusion_matrix(y_test, y_pred)
conf_matrices.append(cm)

fold += 1

print("\n=== Average Cross-Validation Results ===")
print(f"Average Accuracy: {np.mean(accuracies):.4f}")
print(f"Average Precision: {np.mean(precisions):.4f}")
print(f"Average Recall: {np.mean(recalls):.4f}")
print(f"Average F1-score: {np.mean(f1s):.4f}")

=== Average Cross-Validation Results ===
Average Accuracy: 0.7284
Average Precision: 0.7375
Average Recall: 0.7287
Average F1-score: 0.7265

```

More Statistics

```

avg_cm = np.mean(conf_matrices, axis=0).astype(int)
# Print in clean table format
print("\nAverage Confusion Matrix:")
labels = list(label_map.values())
header = "Predicted →\t\t" + "\t".join(labels)
print(header)
for i, row in enumerate(avg_cm):
    row_str = "\t".join(f"{val:.2f}" for val in row)
    print(f"Actual {labels[i]}:\t{row_str}")

```

Average Confusion Matrix:

Predicted →	cats	panda	spiders
Actual cats:	136.00	39.00	24.00
Actual panda:	24.00	169.00	4.00
Actual spiders:	34.00	36.00	130.00

Scale the data

```
# Normalize to improve model performance
scaler = StandardScaler()
X = scaler.fit_transform(X)

import joblib
joblib.dump(scaler, "NB_scaler.pkl")

['NB_scaler.pkl']
```

Add the final module

```
final_model = GaussianNB()
final_model.fit(X, y)
# save the final model:
with open("NB_model.pkl", "wb") as f:
    pickle.dump(final_model, f)
```